W2 - Basic Statistics

Mean - most commonly used average.

Calculation: Sum divided by count

Intuitive: Represents the number you could replace all individual values with to get the same total

Easily skewed by outliers

$$\bar{x} = \frac{\sum x}{N}$$

$\bar{x}$: Mean
$\sum x$: Sum of all values
N : Number of values

Median - middle value.

Calculation:

N is odd, find the middle number when ordered

N is even, find the average of the two middle numbers when ordered

Handles outliers well.

$$\tilde{x} = \left(\frac{N+1}{2}\right)^{th} term$$

$$\tilde{x} = \frac{\left(\frac{N}{2}\right)^{th} term + \left(\frac{N+2}{2}\right)^{th} term}{2}$$

$\tilde{x}$: Median
N : Number of values

Mode - most frequently occurring value.

Other averages

- **Geometric Mean**: percentages, areas, volumes..
- **Harmonic Mean**: used for rates mph.
- **Weighted Average**: a type of mean where some values count for more than others, often used in course gradings.

Standard Deviation - how spread your data is.

Larger SD = More speed

$$\sigma = \sqrt{\frac{\sum(x - \bar{x})^2}{N}}$$

σ : Standard Deviation
Σ: Sum
x : Each value in the dataset
$\bar{x}$: Mean
N: Number of values

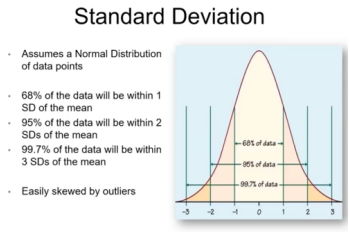

Distributions - how data is distributed.

Histograms - chart that plots the distribution of numeric variables as a series of bars.

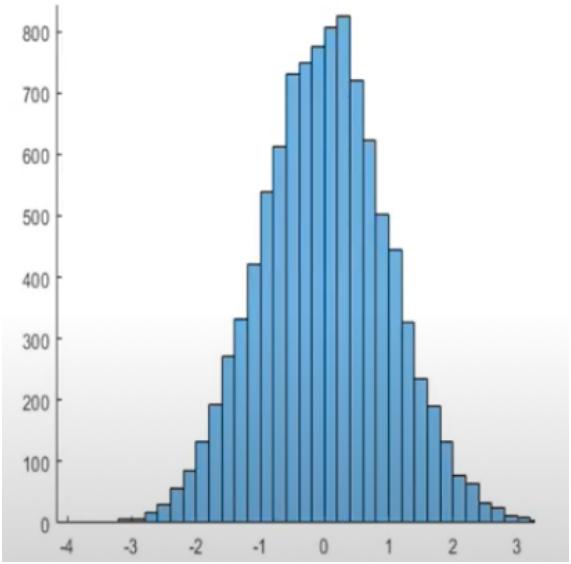

Normal Distribution - Bell curve/ Gaussian Distribution

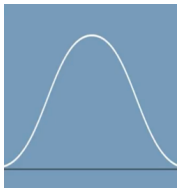

Data here distributed symmetrically with mean and median in the middle.

Kernel Density Estimates - y axis showing probabilities, not counts, x-axis=mean

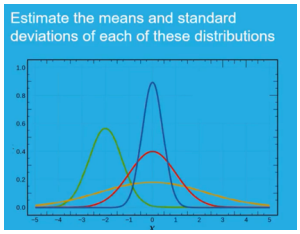

Example:

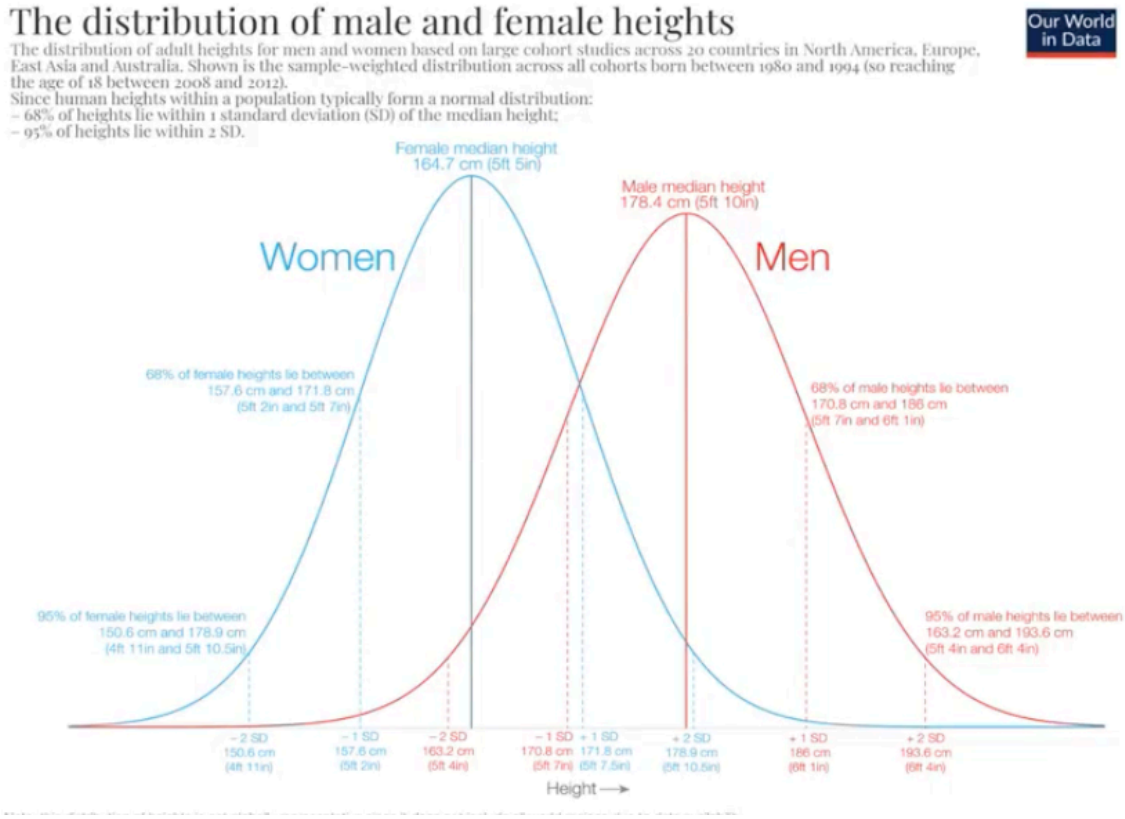

Positive Skew - tail is in the direction of positive, larger values

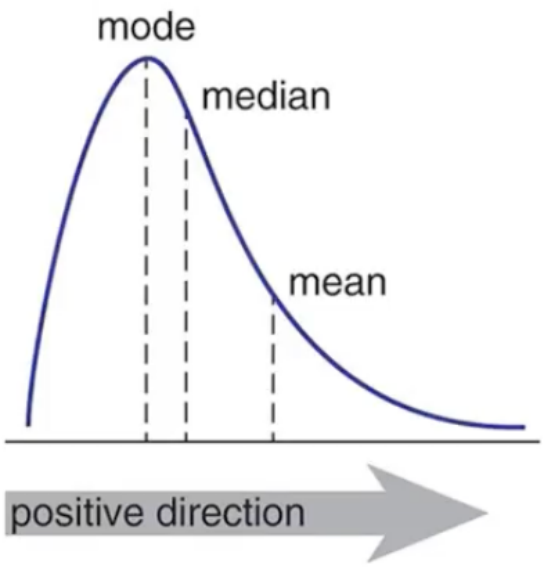

Mean > Median > Mode

Negative Skew - tail is in the direction of negative, smaller values

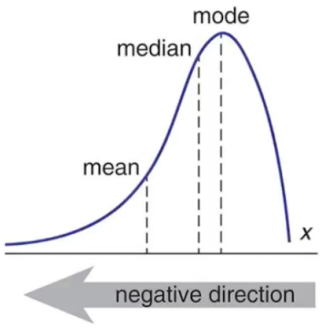

Mean < Median < Mode

Exercise:

- Draw the expected distributions for:
  - Marks achieved by students sitting an A-Level English exam
  - Weights of new-born babies
  - Household incomes in the UK

Characteristics of Normally Distributed Data:

- Symmetrical around the **mean**
- Mean = Median = Mode
- Data within:
  - ±1 standard deviation ≈ 68% of values
  - ±2 standard deviations ≈ 95%
  - ±3 standard deviations ≈ 99.7%

Data that often follow a normal distribution:

- Heights of people in a large population
- Test Scores in large groups
- Daily Temperature Variations

Outliers - a small number of very extreme values can have a great effect on your mean & SD.

What to do with outliers? obvious error or problematic values should be removed.

Correlations - linear relationship between 2 continuous variables

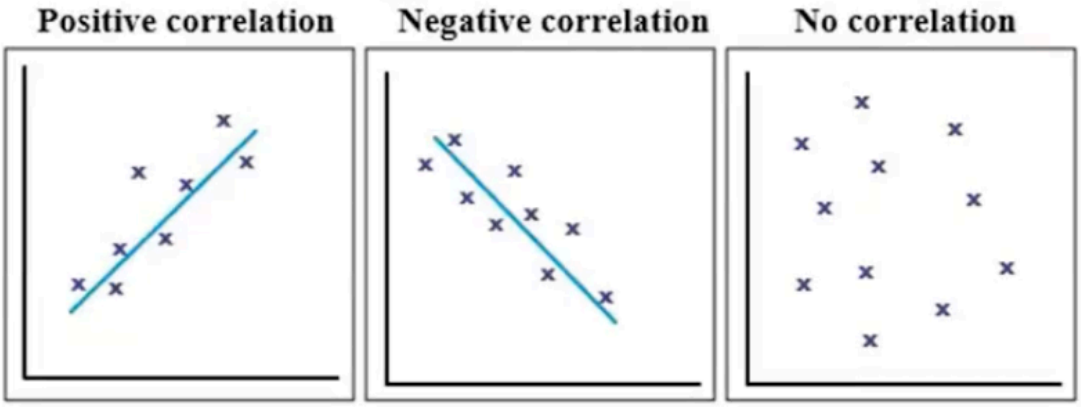

🎯 Think of it like this:

- **Discrete = Counting apples** 🍎: 1, 2, 3, never 2.7 apples.
- **Continuous = Pouring water** 💧: You can measure 1.5L, 1.53L, 1.532L... it never ends.

Examples:

📈Height vs Weight (Taller people tend to weight more)

📚 Study Time vs Test Scores

Correlation Coefficient - r = how closely two variables are related

Pearson's Correlation Coefficient

r = 0 no correlation

r = 1 positive correlation

r = -1 negative correlation

R-Squared R^2 shows how well the data fit the regression model (the goodness of fit) - **scalar(distance) not vector (often used instead of r)**

Calculated based on distance of each point from a line of best fit. Amount of variation in one variable that can be explained by another.

R^2 = 75% means 75% of changes in Y can be accounted for by changes in X

0<=R^2<=1

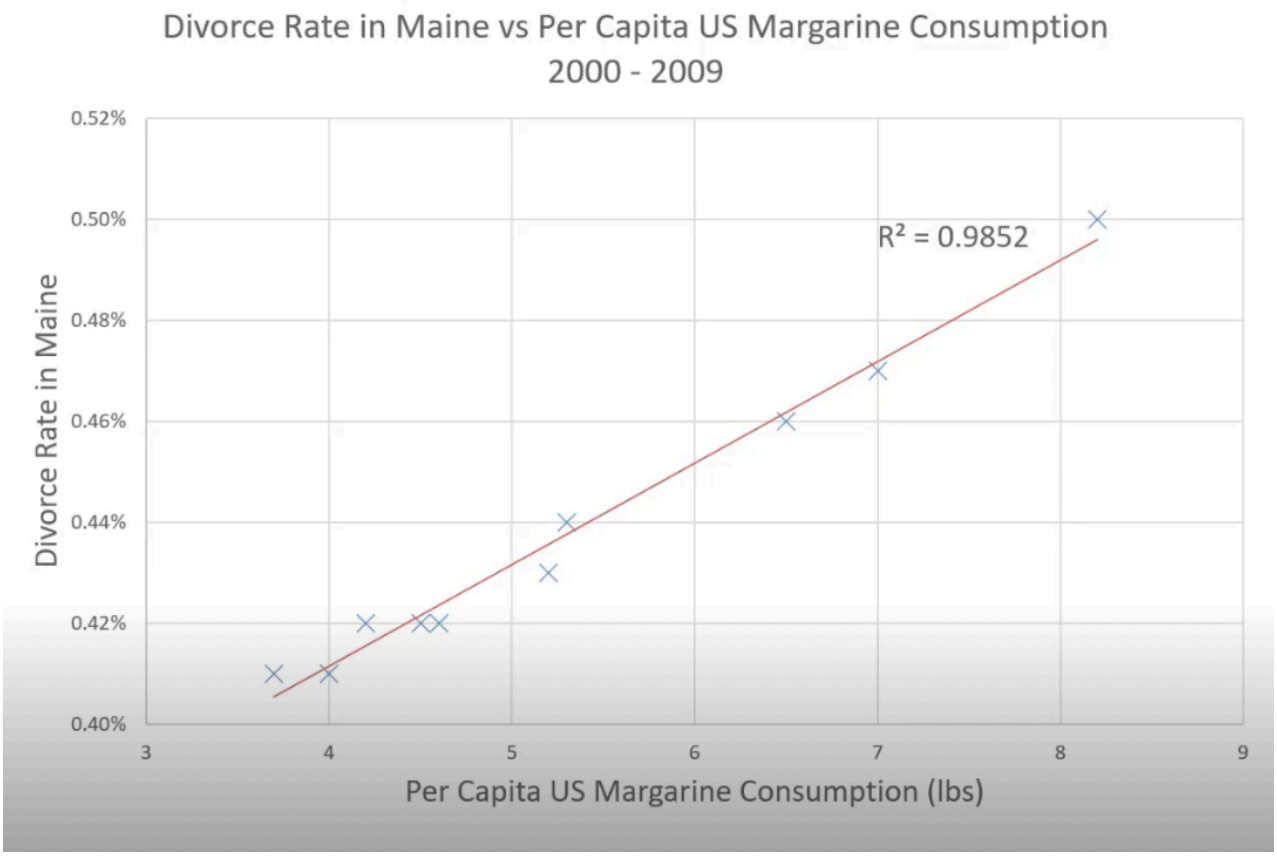

These 2 are correlated even if they're unrelated, so there is a 3rd factor which is almost aways time. No causation at all.